




Learning to Rank

from heuristics to theoretic approaches

Hongning Wang

Congratulations

- Job Offer
 - Design the ranking module for Bing.com



How should I rank documents?

IMAGES VIDEOS MAPS NEWS SEARCH HISTORY MORE | MSN | HOTMAIL

bing™

how to rank documents|



Answer: Rank by relevance!

Relevance ?!

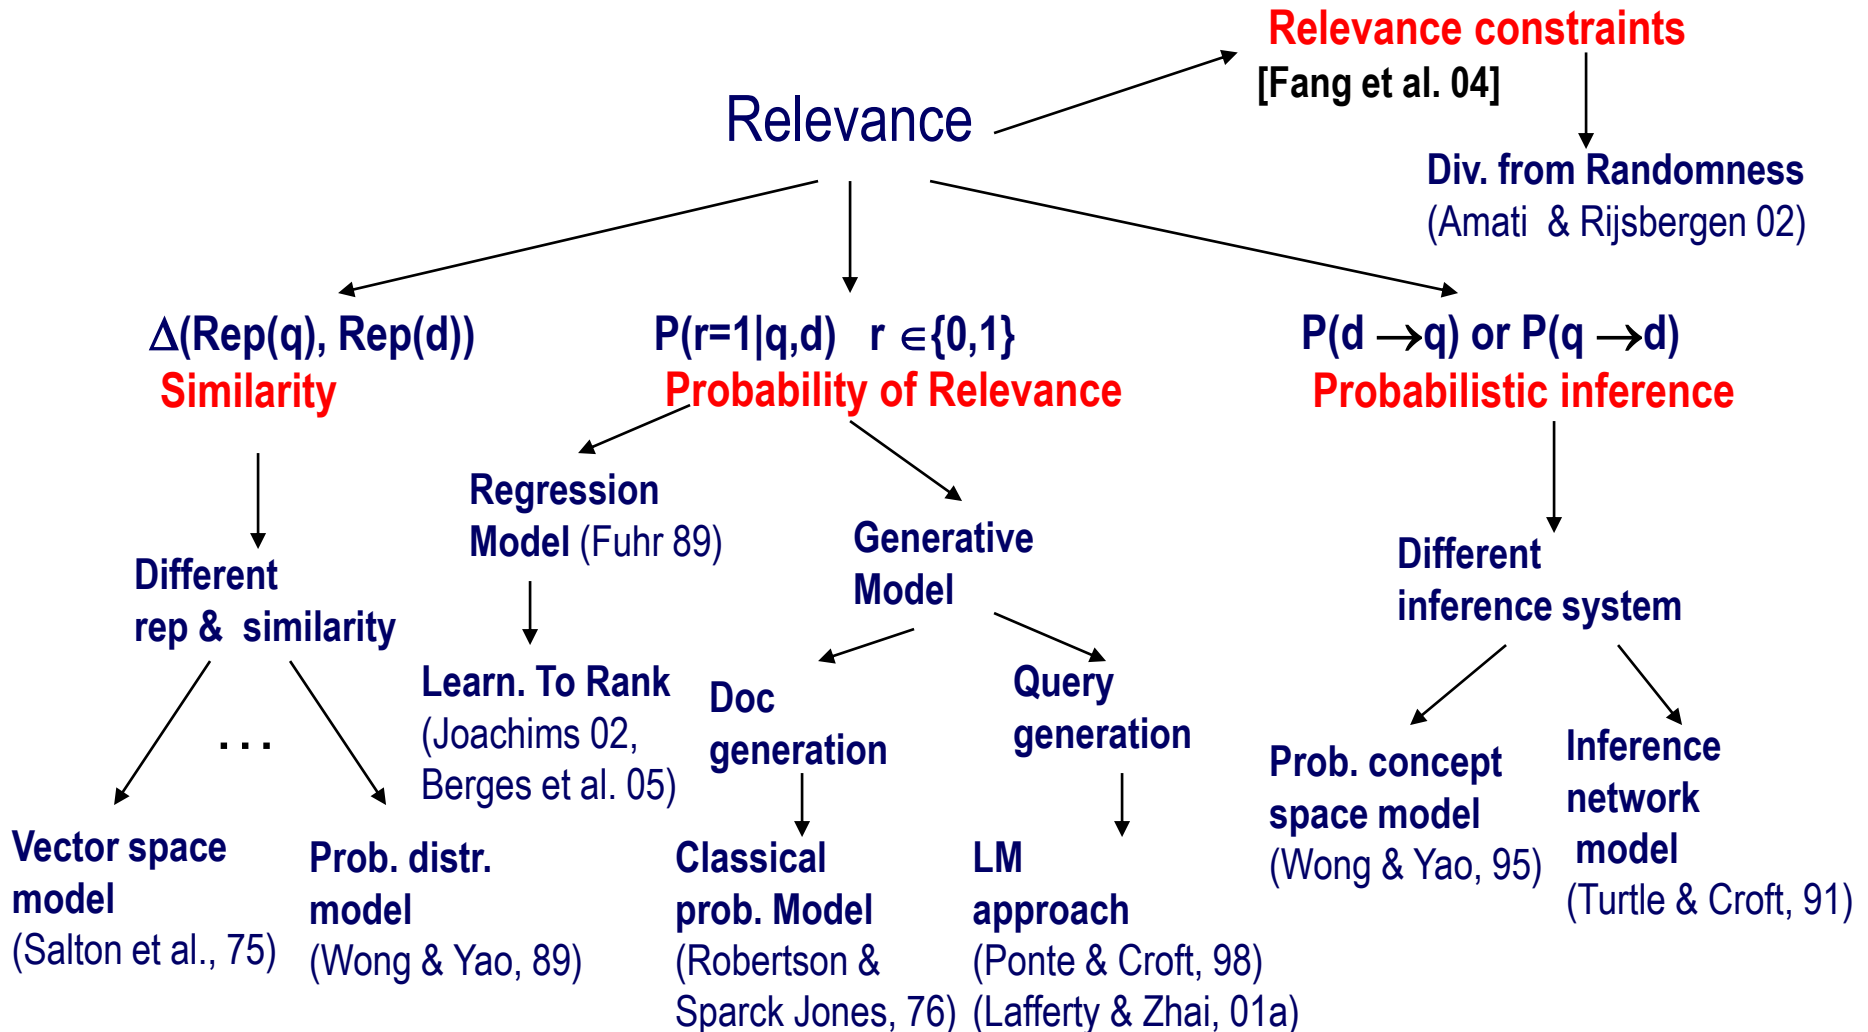
IMAGES VIDEOS MAPS NEWS SEARCH HISTORY MORE | MSN | HOTMAIL

bing™

How to characterize document relevance



The Notion of Relevance



Relevance Estimation

- Query matching
 - Language model
 - BM25
 - Vector space cosine similarity
- Document importance
 - PageRank
 - HITS

Did I do a good job of ranking documents?

- IR evaluations metrics

- Precision

- MAP

- NDCG



how to rank documents

About 128,000,000 results (0.25 seconds)

[Documents as geometric objects: **how to rank documents** for full-text ...](http://www.michaelnielsen.org/.../documents-as-geometric-objects-how-to-...)
www.michaelnielsen.org/.../documents-as-geometric-objects-how-to-...
Jul 7, 2011 – In this post I explain the basic ideas of **how to rank** different **documents** according to their relevance. The ideas used are very beautiful.

[PDF] [Information Retrieval: **Ranking Documents**](http://ciir.cs.umass.edu/~strohman/slides/IR-Intro-Ranking.pdf)

ciir.cs.umass.edu/~strohman/slides/IR-Intro-Ranking.pdf

File Format: PDF/Adobe Acrobat - [View as HTML](#)

Web features, implicit relevance indicators. • Evaluating ranking quality. • Test collections. • Quality metrics. • Training systems to **rank documents** better. 10 ...

[lucene.net - Lucene: **How to rank documents** according to the ...](http://stackoverflow.com/.../lucene-how-to-rank-documents-according-to-t...)
stackoverflow.com/.../lucene-how-to-rank-documents-according-to-t...

1 answer - Mar 3

Top answer: This will require some work, but you can achieve this using payloads. See answers to this very similar question: How to get a better Lucene/Solr score ...

[The Anatomy of a Search Engine](http://infolab.stanford.edu/~backrub/google.html)

infolab.stanford.edu/~backrub/google.html

We use font size relative to the rest of the **document** because when searching, you do not want to **rank** otherwise identical **documents** differently just because ...

CS 6501: Information Retrieval

Take advantage of different relevance estimator?



- Ensemble the cues
 - Linear?

- $$a_1 \times BM25 + \alpha_2 \times LM + \alpha_3 \times PageRank + \alpha_4 \times HITS$$

~~Non-linear?~~

~~Decision tree~~


~~$\{\alpha_1 = 0.1, \alpha_2 = 0.1, \alpha_3 =$~~



$= 0.20, NDCG = 0.6\}$

$= 0.12, NDCG = 0.5\}$


$= 0.18, NDCG = 0.7\}$




What if we have thousands of features?

- Is there any way I can do better?
 - Optimizing the metrics automatically!

Where to find those tree structures?




How to determine those α s?



Rethink the task

- Given: (query, document) pairs represented by a set of relevance estimators, a.k.a., features

DocID	BM25	LM	PageRank	Label
0001	1.6	1.1	0.9	0
0002	2.7	1.9	0.2	1

- Needed: a way of combining the estimators
 - $f(q, \{d\}_{i=1}^D) \rightarrow \text{ordered } \{d\}_{i=1}^D$
- Criterion: optimize IR metrics  **Key!**
 - P@k, MAP, NDCG, etc.

Machine Learning


- Input: $\{(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)\}$, where $X_i \in \mathbb{R}^N, Y_i \in \mathbb{R}^M$
- Object function : $O(Y', Y)$
- Output: $f(X) \rightarrow Y$, such that $f = \operatorname{argmax}_{f' \in F} O(f'(X), Y)$

NOTE: We will only talk about supervised learning.

Classification

$$O(Y', Y) = \delta(Y' = Y)$$


http://en.wikipedia.org/wiki/Statistical_classification



Regression

$$O(Y', Y) = -||Y' - Y||$$

http://en.wikipedia.org/wiki/Regression_analysis



Learning to Rank

- General solution in optimization framework
 - Input: $\{((q_i, d_1), y_1), ((q_i, d_2), y_2), \dots, ((q_i, d_n), y_n)\}$,
where $d_n \in R^N, y_i \in \{0, \dots, L\}$
 - Object: $O = \{P@k, MAP, NDCG\}$
 - Output: $f(q, d) \rightarrow Y$, s.t., $f = \operatorname{argmax}_{f' \in F} O(f'(q, d), Y)$

DocID	BM25	LM	PageRank	Label
0001	1.6	1.1	0.9	0
0002	2.7	1.9	0.2	1

Challenge: how to optimize?

- Evaluation metric recap

- Average Precision

- $$\text{AveP} = \frac{\sum_{k=1}^n (P(k) \times \text{rel}(k))}{\text{number of relevant documents}}$$

- DCG

- $$\text{DCG}_p = \text{rel}_1 + \sum_{i=2}^p \frac{\text{rel}_i}{\log_2 i}$$

Not continuous with respect to $f(X)$!

- Order is essential!

- $f \rightarrow \text{order} \rightarrow \text{metric}$



Approximating the Objects!

- Pointwise
 - Fit the relevance labels individually
- Pairwise
 - Fit the relative orders
- Listwise
 - Fit the whole order



Pointwise Learning to Rank

- Ideally perfect relevance prediction leads to perfect ranking
 - $f \rightarrow \mathbf{score} \rightarrow \text{order} \rightarrow \text{metric}$
- Reducing ranking problem to
 - Regression
 - $O(f(Q, D), Y) = -\sum_i ||f(q_i, d_i) - y_i||$
 - Subset Ranking using Regression, D.Cossock and T.Zhang, COLT 2006
 - (multi-)Classification
 - $O(f(Q, D), Y) = \sum_i \delta(f(q_i, d_i) = y_i)$
 - Ranking with Large Margin Principles, A. Shashua and A. Levin, NIPS 2002

Subset Ranking using Regression

D.Cossock and T.Zhang, COLT 2006

- Fit relevance labels via regression


- $\hat{f} = \arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \left[\sum_{j=1}^m (f(x_{i,j}, S_i) - y_{i,j})^2 \right]$

- Emphasize more on relevant documents

- $\sum_{j=1}^m w(x_j, S)(f(x_j, S) - y_j)^2 + u \sup_j w'(x_j, S)(f(x_j, S) - \delta(x_j, S))_+^2$

Weights on each document


Most positive document



Ranking with Large Margin Principles

A. Shashua and A. Levin, NIPS 2002

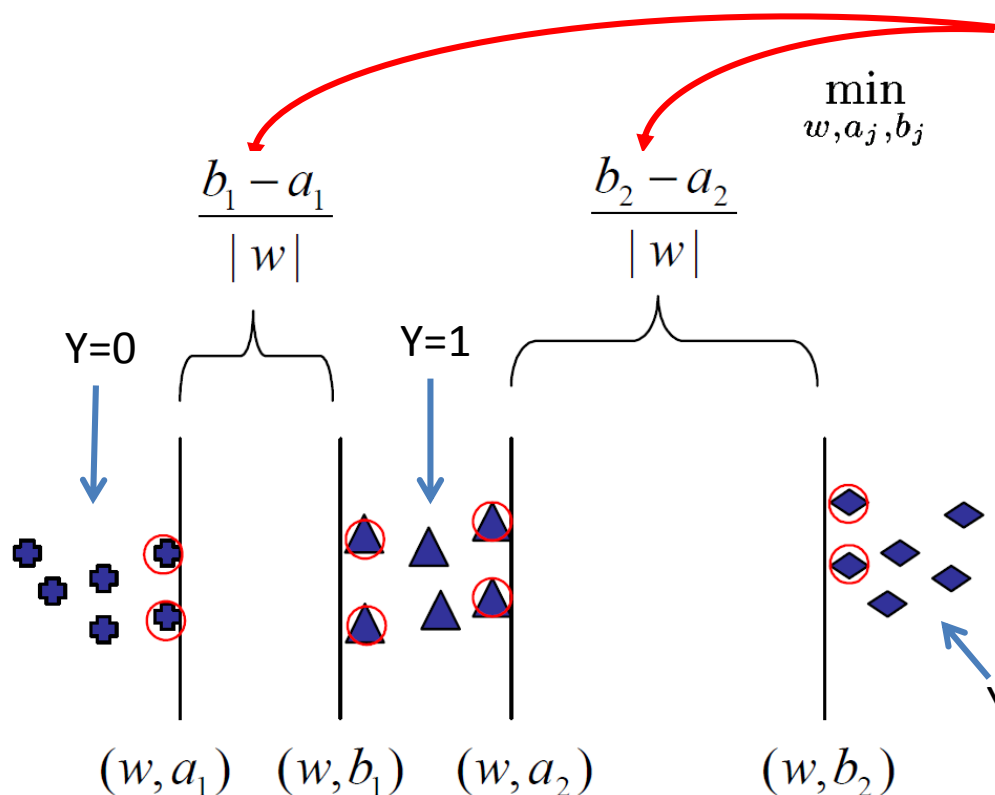
- Goal: correctly placing the documents in the corresponding category and maximize the margin



Ranking with Large Margin Principles

A. Shashua and A. Levin, NIPS 2002

- Maximizing the sum of margins



$$\min_{w, a_j, b_j} \sum_{j=1}^{k-1} (a_j - b_j) + C \sum_i \sum_j (\epsilon_i^j + \epsilon_i^{*j+1})$$

subject to

$$a_j \leq b_j,$$

$$b_j \leq a_{j+1}, \quad j = 1, \dots, k-2$$

$$\mathbf{w} \cdot \mathbf{x}_i^j \leq a_j + \epsilon_i^j, \quad b_j - \epsilon_i^{*j+1} \leq \mathbf{w} \cdot \mathbf{x}_i^{j+1}$$


$$\mathbf{w} \cdot \mathbf{w} \leq 1, \quad \epsilon_i^j \geq 0, \epsilon_i^{*j+1} \geq 0$$

Ranking with Large Margin Principles

A. Shashua and A. Levin, NIPS 2002

- Ranking loss is consistently decreasing with more training data

Crammer & Singer 2001



What did we learn



- Machine learning helps!
 - Derive something optimizable
 - More efficient and guided

There is always a catch

- Cannot directly optimize IR metrics
 - $(0 \rightarrow 1, 2 \rightarrow 0)$ worse than $(0 \rightarrow -2, 2 \rightarrow 4)$
- Position of documents are ignored
 - Penalty on documents at higher positions should be larger
- Favor the queries with more documents

Pairwise Learning to Rank

- Ideally perfect partial order leads to perfect ranking
 - $f \rightarrow$ **partial order** \rightarrow order \rightarrow metric
- Ordinal regression
 - $O(f(Q, D), Y) = \sum_{i \neq j} \delta(y_i > y_j) \delta(f(q_i, d_i) > f(q_j, d_j))$
 - Relative ordering between different documents is significant
 - E.g., (0 \rightarrow -2, 2 \rightarrow 4) is better than (0 \rightarrow 1, 2 \rightarrow 0)
 - Large body of work

Optimizing Search Engines using Clickthrough Data

Thorsten Joachims, KDD'02

- Minimizing the number of mis-ordered pairs

$y_1 > y_2, y_2 > y_3, y_1 > y_4$
 $y_1 > y_2$
linear combination of features
 $f(q, d) = w^T X_{q,d}$

$$\text{minimize:} \quad V(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w} \cdot \vec{w} + C \sum \xi_{i,j,k}$$

subject to:

$$\forall (d_i, d_j) \in r_1^* : \vec{w}\Phi(q_1, d_i) \geq \vec{w}\Phi(q_1, d_j) + 1 - \xi_{i,j,1}$$


...

$$\forall (d_i, d_j) \in r_n^* : \vec{w}\Phi(q_n, d_i) \geq \vec{w}\Phi(q_n, d_j) + 1 - \xi_{i,j}, \mathbf{1}$$

$$\forall i \forall j \forall k : \xi_{i,j,k} \geq 0$$

Keep the relative orders


RankingSVM



Optimizing Search Engines using Clickthrough Data

Thorsten Joachims, KDD'02

- How to use it?
 - $f \rightarrow \mathbf{score} \rightarrow \mathbf{order}$



Optimizing Search Engines using Clickthrough Data

Thorsten Joachims, KDD'02

- What did it learn from the data?
 - Linear correlations

Positive
correlated
features



Negative
correlated
features




weight	feature
0.60	query_abstract_cosine
0.48	top10_google
0.24	query_url_cosine
0.24	top1count_1
0.24	top10_msnsearch
0.22	host_citeseer
0.21	domain_nec
0.19	top10count_3
0.17	top1_google
0.17	country_de
...	
0.16	abstract_contains_home
0.16	top1_hotbot
...	
0.14	domain_name_in_query
...	
-0.13	domain_tu-bs
-0.15	country_fi
-0.16	top50count_4
-0.17	url_length
-0.22	top10count_0
-0.32	top1count_0
-0.38	top1count_0

Optimizing Search Engines using Clickthrough Data

Thorsten Joachims, KDD'02


- How good is it?
 - Test on real system



An Efficient Boosting Algorithm for Combining Preferences

Y. Freund, R. Iyer, et al. JMLR 2003

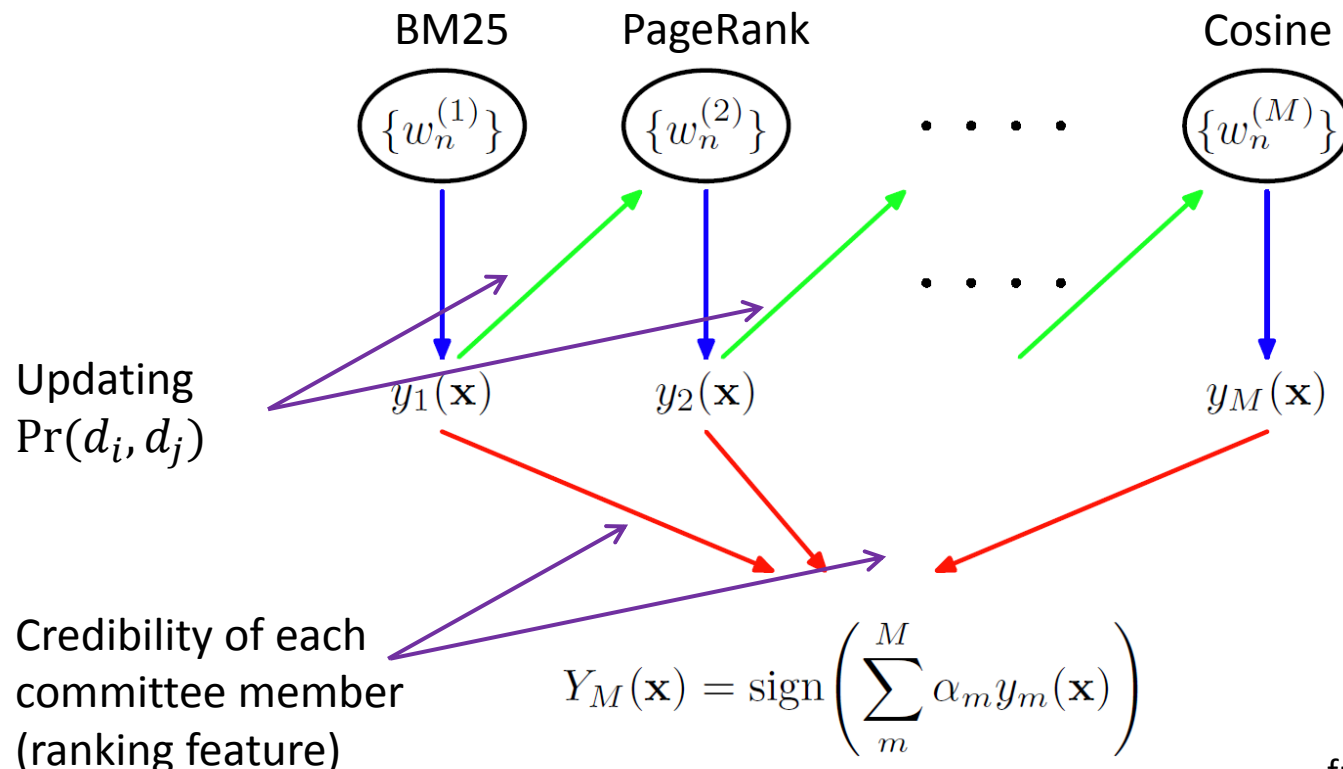
- Smooth the loss on mis-ordered pairs
$$-\sum_{y_i > y_j} Pr(d_i, d_j) \underline{\exp[f(q, d_j) - f(q, d_i)]}$$



An Efficient Boosting Algorithm for Combining Preferences

Y. Freund, R. Iyer, et al. JMLR 2003


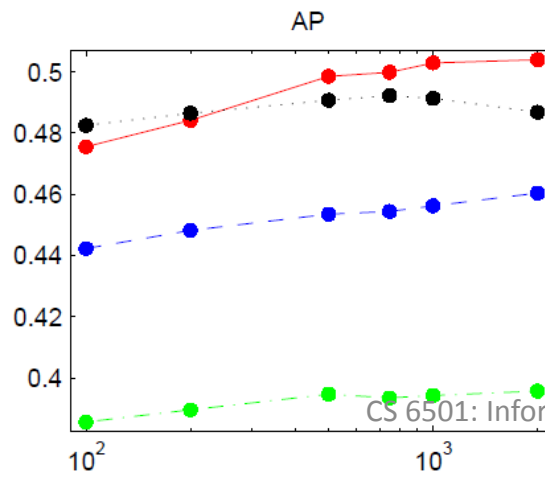
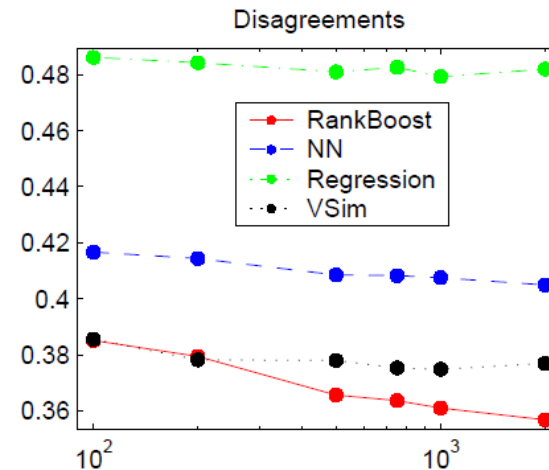

- RankBoost: optimize via boosting
 - Vote by a committee



An Efficient Boosting Algorithm for Combining Preferences

Y. Freund, R. Iyer, et al. JMLR 2003

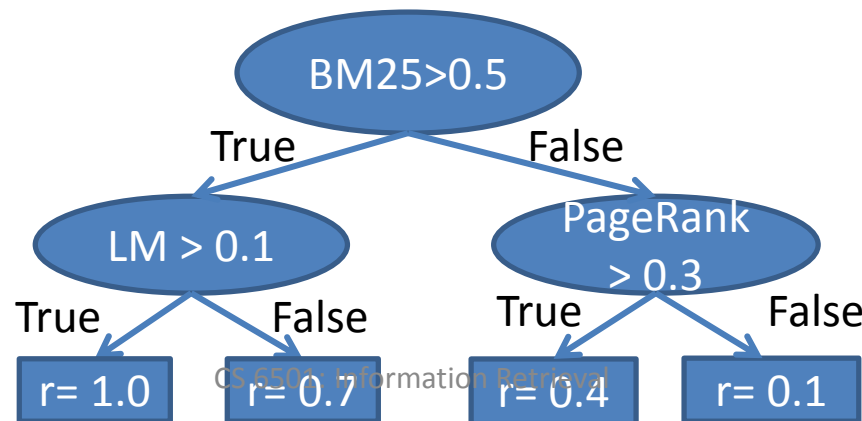
- How good is it?



A Regression Framework for Learning Ranking Functions Using Relative Relevance Judgments

Zheng et al. SIRIG'07


- Non-linear ensemble of features
 - Object: $\sum_{y_i > y_j} (\max\{0, f(q, d_j) - f(q, d_i)\})^2$
 - Gradient descent boosting tree
 - Boosting tree
 - Using regression tree to minimize the residuals
 - $r^t(q, d, y) = O^t(q, d, y) - f^{(t-1)}(q, d, y)$



A Regression Framework for Learning Ranking Functions Using Relative Relevance Judgments

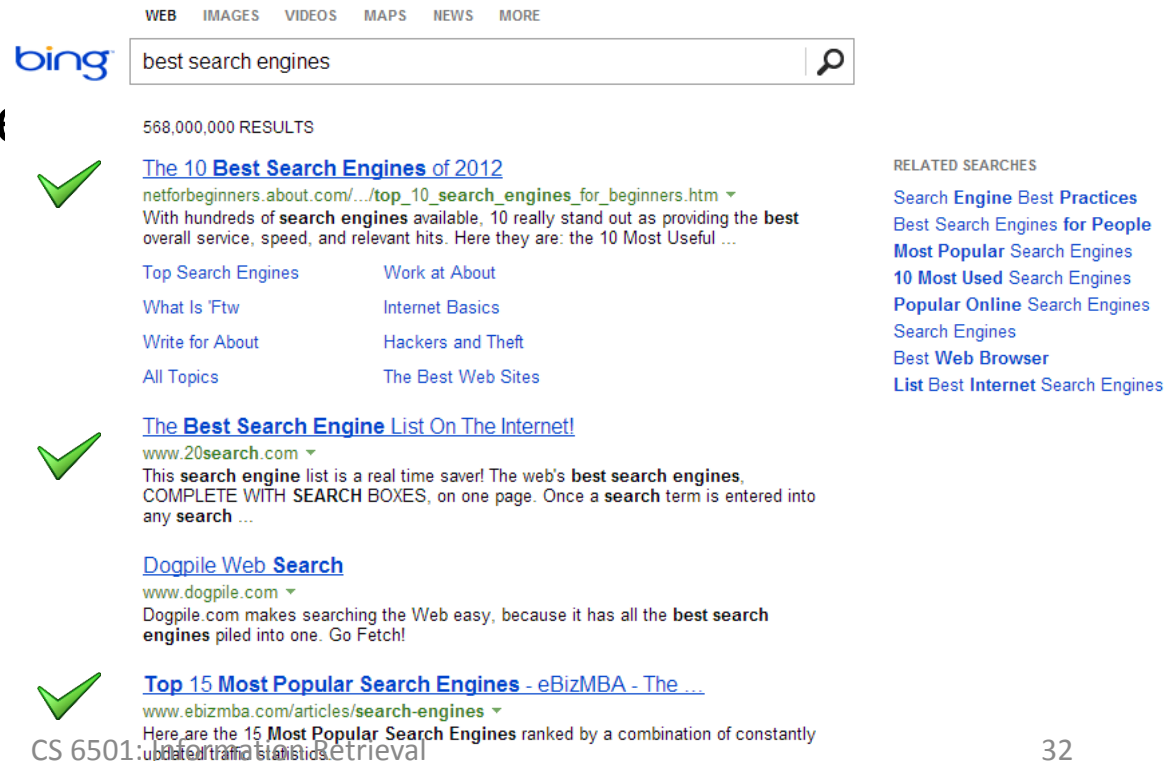
Zheng et al. SIRIG'07

- Non-linear v.s. linear
 - Comparing with RankingSVM



Where do we get the relative orders

- Human annotations
 - Small scale, expensive to acquire
- Clickthroughs
 - Large amount, (



The screenshot shows a Bing search results page for the query "best search engines". The search bar at the top contains the text "best search engines" and a magnifying glass icon. Below the search bar, it says "568,000,000 RESULTS". The first three search results are highlighted with green checkmarks on the left side of the slide. The first result is titled "The 10 Best Search Engines of 2012" and is from the website "netforbeginners.about.com". The second result is titled "The Best Search Engine List On The Internet!" and is from "www.20search.com". The third result is titled "Dogpile Web Search" and is from "www.dogpile.com". To the right of the search results, there is a section titled "RELATED SEARCHES" with several links: "Search Engine Best Practices", "Best Search Engines for People", "Most Popular Search Engines", "10 Most Used Search Engines", "Popular Online Search Engines", "Search Engines", "Best Web Browser", and "List Best Internet Search Engines".

WEB IMAGES VIDEOS MAPS NEWS MORE

bing best search engines

568,000,000 RESULTS

The 10 Best Search Engines of 2012
netforbeginners.about.com/.../top_10_search_engines_for_beginners.htm
With hundreds of search engines available, 10 really stand out as providing the best overall service, speed, and relevant hits. Here they are: the 10 Most Useful ...

Top Search Engines Work at About
What Is 'Ftw Internet Basics
Write for About Hackers and Theft
All Topics The Best Web Sites

The Best Search Engine List On The Internet!
www.20search.com
This search engine list is a real time saver! The web's best search engines, COMPLETE WITH SEARCH BOXES, on one page. Once a search term is entered into any search ...

Dogpile Web Search
www.dogpile.com
Dogpile.com makes searching the Web easy, because it has all the best search engines piled into one. Go Fetch!

Top 15 Most Popular Search Engines - eBizMBA - The ...
www.ebizmba.com/articles/search-engines
Here are the 15 Most Popular Search Engines ranked by a combination of constantly updated traffic statistics.

RELATED SEARCHES

Search Engine Best Practices
Best Search Engines for People
Most Popular Search Engines
10 Most Used Search Engines
Popular Online Search Engines
Search Engines
Best Web Browser
List Best Internet Search Engines

Accurately Interpreting Clickthrough Data as Implicit Feedback

Thorsten Joachims, et al., SIGIR'05

- Position bias

Your click is not because

Table 2: Percentage of times the user viewed an abstract at a particular rank before he clicked on a link at a particular rank.

Viewed Rank	Clicked Rank					
	1	2	3	4	5	6
1	90.6%	76.2%	73.9%	60.0%	54.5%	45.5%
2	56.8%	90.5%	82.6%	53.3%	63.6%	54.5%
3	30.2%	47.6%	95.7%	80.0%	81.8%	45.5%
4	17.3%	19.0%	47.8%	93.3%	63.6%	45.5%
5	8.6%	14.3%	21.7%	53.3%	100.0%	72.7%
6	4.3%	4.8%	8.7%	33.3%	18.2%	81.8%

Accurately Interpreting Clickthrough Data as Implicit Feedback

Thorsten Joachims, et al., SIGIR'05

- Controlled experiment
 - Over trust of the top ranked positions

“normal”	l_1^-, l_2^-	l_1^+, l_2^-	l_1^-, l_2^+	l_1^+, l_2^+	total
$\text{rel}(l_1) > \text{rel}(l_2)$	15	19	1	1	36
$\text{rel}(l_1) < \text{rel}(l_2)$	11	5	2	2	20
$\text{rel}(l_1) = \text{rel}(l_2)$	19	9	1	0	29
total	45	33	4	3	85

“swapped”	l_1^-, l_2^-	l_1^+, l_2^-	l_1^-, l_2^+	l_1^+, l_2^+	total
$\text{rel}(l_1) > \text{rel}(l_2)$	11	15	1	1	28
$\text{rel}(l_1) < \text{rel}(l_2)$	17	10	7	2	36
$\text{rel}(l_1) = \text{rel}(l_2)$	36	11	3	0	50
total	64	36	11	3	114

Accurately Interpreting Clickthrough Data as Implicit Feedback

Thorsten Joachims, et al., SIGIR'05

- Pairwise preference matters
 - Click: examined and clicked document
 - Skip: examined but non-clicked document

Explicit Feedback Data Strategy	Abstracts					Pages Phase II all
	Phase I “normal”	“normal”	Phase II “swapped”	“reversed”	all	
Inter-Judge Agreement	89.5	N/A	N/A	N/A	82.5	86.4
Click > Skip Above	80.8 ± 3.6	88.0 ± 9.5	79.6 ± 8.9	83.0 ± 6.7	83.1 ± 4.4	78.2 ± 5.6
Last Click > Skip Above	83.1 ± 3.8	89.7 ± 9.8	77.9 ± 9.9	84.6 ± 6.9	83.8 ± 4.6	80.9 ± 5.1
Click > Earlier Click	67.2 ± 12.3	75.0 ± 25.8	36.8 ± 22.9	28.6 ± 27.5	46.9 ± 13.9	64.3 ± 15.4
Click > Skip Previous	82.3 ± 7.3	88.9 ± 24.1	80.0 ± 18.0	79.5 ± 15.4	81.6 ± 9.5	80.7 ± 9.6
Click > No Click Next	84.1 ± 4.9	75.6 ± 14.5	66.7 ± 13.1	70.0 ± 15.7	70.4 ± 8.0	67.4 ± 8.2

Click > Skip

What did we learn



- Predicting relative order
 - Getting closer to the nature of ranking
- Promising performance in practice
 - Pairwise preferences from click-throughs

Listwise Learning to Rank

- Can we directly optimize the ranking?
 - $f \rightarrow \mathbf{order} \rightarrow \text{metric}$
- Tackle the challenge
 - Optimization without gradient



From RankNet to LambdaRank to LambdaMART: An Overview


Christopher J.C. Burges, 2010

- Minimizing mis-ordered pair => maximizing IR metrics?

Mis-ordered pairs: 6

$$AP: \frac{5}{8}$$

DCG: 1.333



Mis-ordered pairs: 4

$$AP: \frac{5}{12}$$

DCG: 0.931

Position is crucial!

From RankNet to LambdaRank to LambdaMART: An Overview

Christopher J.C. Burges, 2010

- Weight the mis ordered pairs?

– Score

right

– In




– In

the

document i, j

ΔG , if
saving



Gradient

i.e., export

CS@UVA

From RankNet to LambdaRank to LambdaMART: An Overview

Christopher J.C. Burges, 2010

- Lambda functions
 - Gradient?
 - Yes, it meets the sufficient and necessary condition of being partial derivative
 - Lead to optimal solution of original problem?
 - Empirically

From RankNet to LambdaRank to LambdaMART: An Overview

Christopher J.C. Burges, 2010

- Evolution

	RankNet
Object	Cross entropy over the pairs
Gradient (λ function)	Gradient of cross entropy
Optimization method	neural network



As we discussed
in RankBoost

Optimize solely
by gradient

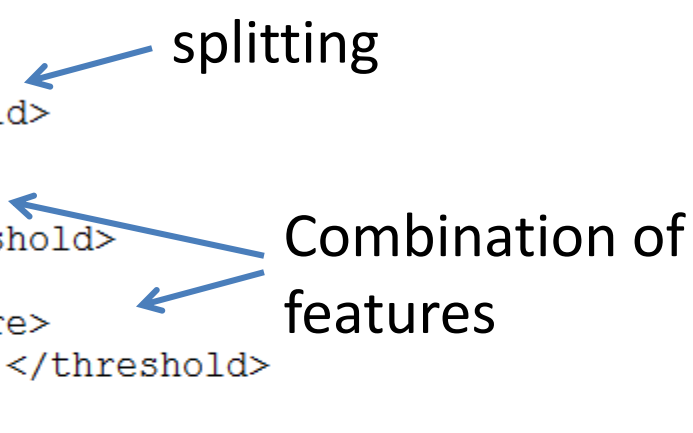
Non-linear
combination

From RankNet to LambdaRank to LambdaMART: An Overview

Christopher J.C. Burges, 2010

- A Lambda tree

```
<tree id="8" weight="0.1">
  <split>
    <feature> 811 </feature>
    <threshold> 5.0 </threshold>
    <split pos="left">
      <feature> 33 </feature>
      <threshold> 20.0 </threshold>
      <split pos="left">
        <feature> 589 </feature>
        <threshold> 43493.125 </threshold>
        <split pos="left">
          <feature> 1094 </feature>
          <threshold> 302.73438 </threshold>
          <split pos="left">
            <feature> 108 </feature>
            <threshold> 9881.824 </threshold>
            <split pos="left">
              <output> -0.66917753 </output>
            </split>
            <split pos="right">
              <feature> 151 </feature>
              <threshold> 9072276.0 </threshold>
```



The diagram illustrates the structure of a Lambda tree. It shows a series of nested split operations. The root split is on feature 811 with threshold 5.0. The left child of this split is another split on feature 33 with threshold 20.0. This pattern continues with splits on features 589, 1094, and 108. The final split shown is on feature 108 with threshold 9881.824, which leads to an output of -0.66917753. The right child of the final split is a split on feature 151 with threshold 9072276.0. Annotations with arrows point to specific parts of the tree: 'splitting' points to the root split, and 'Combination of features' points to the nested structure of splits.

AdaRank: a boosting algorithm for information retrieval

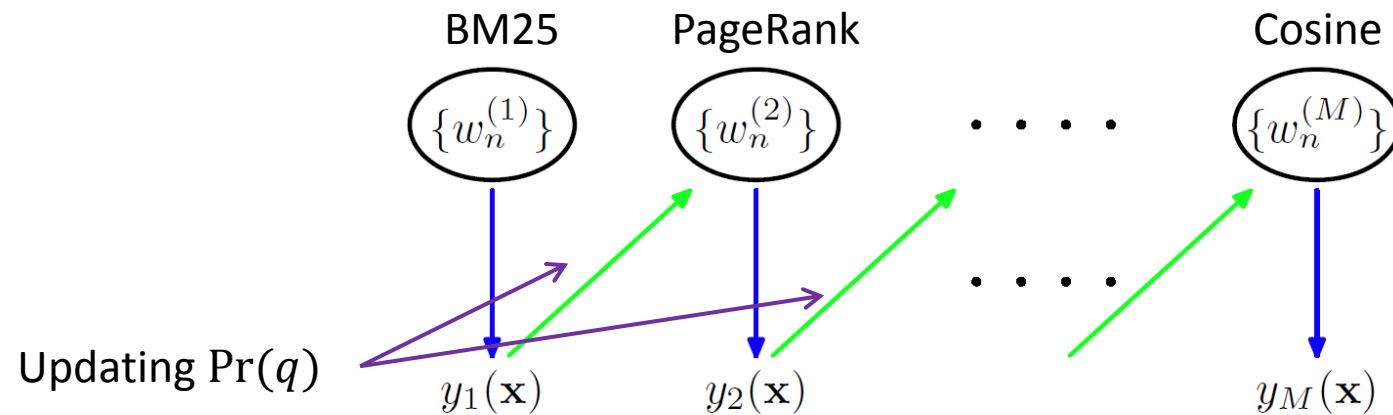
Jun Xu & Hang Li, SIGIR'07

- Loss defined by IR metrics

- $-\sum_{q \in Q} Pr(q) \exp[-O(q)]$

Target metrics: MAP, NDCG, MRR

- Optimizing by boosting



Credibility of each committee member (ranking feature)

$$Y_M(\mathbf{x}) = \text{sign} \left(\sum_{m=1}^M \alpha_m y_m(\mathbf{x}) \right)$$

A Support Vector Machine for Optimizing Average Precision

Yisong Yue, et al., SIGIR'07

RankingSVM

- Minimizing the pairwise loss

$$\begin{aligned}
 &\text{minimize:} \quad V(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w} \cdot \vec{w} + C \sum \xi_{i,j,k} \\
 &\text{subject to:} \\
 &\quad \underline{\forall (d_i, d_j) \in r_1^* : \vec{w} \Phi(q_1, d_i) \geq \vec{w} \Phi(q_1, d_j) + 1 - \xi_{i,j,1}} \\
 &\quad \dots \\
 &\quad \forall (d_i, d_j) \in r_n^* : \vec{w} \Phi(q_n, d_i) \geq \vec{w} \Phi(q_n, d_j) + 1 - \xi_{i,j,n} \\
 &\quad \forall i \forall j \forall k : \xi_{i,j,k} \geq 0
 \end{aligned}$$

Loss defined on the number of mis-ordered document pairs

SVM-MAP

- Minimizing the structural loss

$$\begin{aligned}
 &\min_{\mathbf{w}, \xi \geq 0} \frac{1}{2} \|\mathbf{w}\|^2 + \frac{C}{n} \sum_{i=1}^n \xi_i \\
 &s.t. \quad \forall i, \underline{\forall \mathbf{y} \in \mathcal{Y} \setminus \mathbf{y}_i :} \\
 &\quad \mathbf{w}^T \Psi(\mathbf{x}_i, \mathbf{y}_i) \geq \mathbf{w}^T \Psi(\mathbf{x}_i, \mathbf{y}) + \Delta(\mathbf{y}_i, \mathbf{y}) - \xi_i
 \end{aligned}$$


MAP difference

Loss defined on the quality of the whole list of ordered documents

A Support Vector Machine for Optimizing Average Precision

Yisong Yue, et al., SIGIR'07

- Max margin principle
 - Push the ground-truth far away from any mistakes you might make
 - Finding the most violated constraints



A Support Vector Machine for Optimizing Average Precision

Yisong Yue, et al., SIGIR'07


- Finding the most violated constraints
 - MAP is invariant to permutation of (ir)relevant documents
 - Maximize MAP over a series of swaps between relevant and irrelevant documents

- $\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}} \Delta(\mathbf{y}_i, \mathbf{y}) + \mathbf{w}^T \Psi(\mathbf{x}_i, \mathbf{y})$

Right-hand side of constraints

Start from the reverse order of ideal ranking

Greedy solution



A Support Vector Machine for Optimizing Average Precision

Yisong Yue, et al., SIGIR'07

- Experiment results

Model	TREC 9		TREC 10	
	MAP	W/L	MAP	W/L
SVM_{map}^{Δ}	0.290	—	0.287	—
SVM_{roc}^{Δ}	0.282	29/21	0.278	35/15 **
SVM_{acc}	0.213	49/1 **	0.222	49/1 **
SVM_{acc2}	0.270	34/16 **	0.261	42/8 **
SVM_{acc3}	0.133	50/0 **	0.182	46/4 **
SVM_{acc4}	0.233	47/3 **	0.238	46/4 **

Other listwise solutions

- Soften the metrics to make them differentiable
 - Michael Taylor et al., SoftRank: optimizing non-smooth rank metrics, WSDM'08
- Minimize a loss function defined on permutations
 - Zhe Cao et al., Learning to rank: from pairwise approach to listwise approach, ICML'07

What did we learn



- Taking a list of documents as a whole
 - Positions are visible for the learning algorithm
 - Directly optimizing the target metric
- Limitation
 - The search space is huge!

Summary



- Learning to rank
 - Automatic combination of ranking features for optimizing IR evaluation metrics
- Approaches
 - Pointwise
 - Fit the relevance labels individually
 - Pairwise
 - Fit the relative orders
 - Listwise
 - Fit the whole order

Experimental Comparisons

- Ranking performance

Table 7.5 Results on the TD2003 dataset

Algorithm	N@1	N@3	N@10	P@1	P@3	P@10	MAP
Regression	0.320	0.307	0.326	0.320	0.260	0.178	0.241
RankSVM	0.320	0.344	0.346	0.320	0.293	0.188	0.263
RankBoost	0.280	0.325	0.312	0.280	0.280	0.170	0.227
FRank	0.300	0.267	0.269	0.300	0.233	0.152	0.203
ListNet	0.400	0.337	0.348	0.400	0.293	0.200	0.275
AdaRank	0.260	0.307	0.306	0.260	0.260	0.158	0.228
SVM ^{map}	0.320	0.320	0.328	0.320	0.253	0.170	0.245

Experimental Comparisons

- Winning count
 - Over seven different data sets

Table 7.12 Winner Number of Each Algorithm

Algorithm	N@1	N@3	N@10	P@1	P@3	P@10	MAP
Regression	4	4	4	5	5	5	4
RankSVM	21	22	22	21	22	22	24
RankBoost	18	22	22	17	22	23	19
FRank	18	19	18	18	17	23	15
ListNet	29	31	33	30	32	35	33
AdaRank	26	25	26	23	22	16	27
SVM ^{map}	23	24	22	25	20	17	25


Experimental Comparisons

- My experiments
 - 1.2k queries, 45.5K documents with 1890 features
 - 800 queries for training, 400 queries for testing

	MAP	P@1	ERR	MRR	NDCG@5
ListNET	<i>0.2863</i>	<i>0.2074</i>	<i>0.1661</i>	<i>0.3714</i>	<i>0.2949</i>
LambdaMART	0.4644	0.4630	0.2654	0.6105	0.5236
RankNET	0.3005	0.2222	0.1873	0.3816	0.3386
RankBoost	0.4548	0.4370	0.2463	0.5829	0.4866
RankingSVM	0.3507	0.2370	0.1895	0.4154	0.3585
AdaRank	0.4321	0.4111	0.2307	0.5482	0.4421
pLogistic	0.4519	0.3926	0.2489	0.5535	0.4945
Logistic	0.4348	0.3778	0.2410	0.5526	0.4762

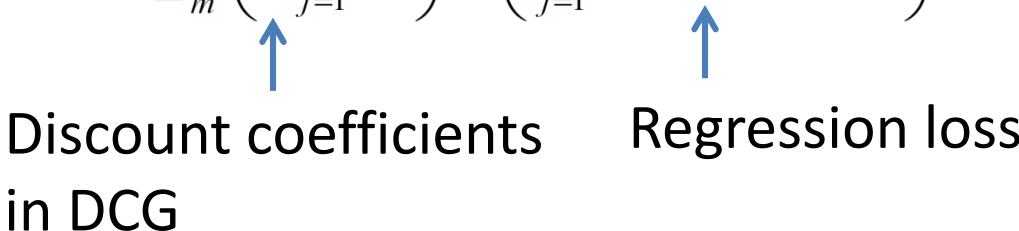
Analysis of the Approaches

- What are they really optimizing?
 - Relation with IR metrics



Pointwise Approaches

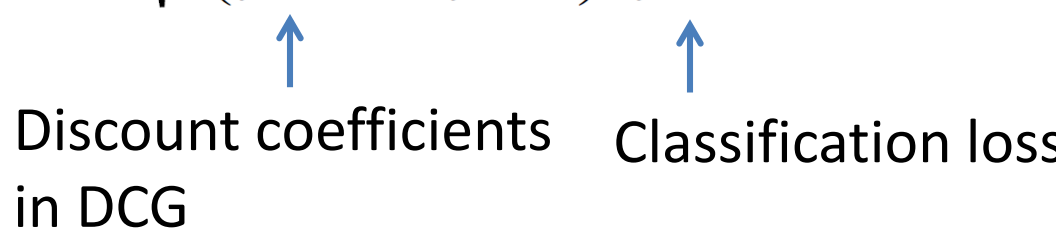
- Regression based

$$1 - NDCG(f) \leq \frac{1}{Z_m} \left(2 \sum_{j=1}^m \eta_j^\varepsilon \right)^{1/\alpha} \left(\sum_{j=1}^m (f(x_j) - y_j)^\beta \right)^{1/\beta}$$


Discount coefficients
in DCG

Regression loss

- Classification based

$$1 - NDCG(f) \leq \frac{15}{Z_m} \sqrt{2 \left(\sum_{j=1}^m \eta_j^2 - m \prod_{j=1}^m \eta_j^{\frac{2}{m}} \right) \cdot \sum_{j=1}^m I_{\{y_j \neq f(x_j)\}}}$$


Discount coefficients
in DCG

Classification loss

Pointwise Approach

- Although it seems the loss functions can bound (1-NDCG), the constants before the losses seem too large.

$$\begin{array}{ccc}
 x_i, y_i & \xrightarrow{\quad} & Z_m \approx 21.4 \\
 \left(\begin{array}{c} x_1, 4 \\ x_2, 3 \\ x_3, 2 \\ x_4, 1 \end{array} \right) & & \left(\begin{array}{c} x_1, 3 \\ x_2, 2 \\ x_3, 1 \\ x_4, 0 \end{array} \right) \\
 & \begin{array}{c} DCG(f) \approx 21.4 \\ \downarrow \\ |1 - NDCG(f)| = 0 \end{array} &
 \end{array}$$

$$\frac{15}{Z_m} \sqrt{2 \left(\sum_{j=1}^m \left(\frac{1}{\log(j+1)} \right)^2 - m \sum_{j=1}^m \left(\frac{1}{\log(j+1)} \right)^{\frac{2}{m}} \right)} \cdot \sum_{j=1}^m I_{\{y_j \neq f(x_j)\}} \approx 1.15 > 1$$

Pairwise Approach

(W. Chen, T.-Y. Liu, et al. 2009)

- Unified loss vs. (1-NDCG) Discount coefficients in DCG
 - When $\beta_t = \frac{G(t)\eta(t)}{Z_m}$, $\tilde{L}(f)$ is a tight bound of (1-NDCG).
- Surrogate function of Unified loss
 - After introducing weights β_t , loss functions in Ranking SVM, RankBoost, RankNet are *Cost-sensitive Pairwise Comparison* surrogate functions, and thus are *consistent* with and are *upper bounds* of the unified loss.
 - Consequently, they also upper bound (1-NDCG).


Listwise Approaches

- No general analysis
 - Method dependent
 - Directness and consistency

Connection with Traditional IR

- People have foreseen this topic long time ago
 - Nicely fit in the risk minimization framework

Applying Bayesian Decision Theory



Metric to be optimized

Available ranking features

$$(D^*, \pi^*) = \arg \min_{D, \pi} \int_{\Theta} \underbrace{L(D, \pi, \theta)}_{\text{loss}} \underbrace{p(\theta | q, U, C, S)}_{\text{hidden observed}} d\theta$$

RISK MINIMIZATION

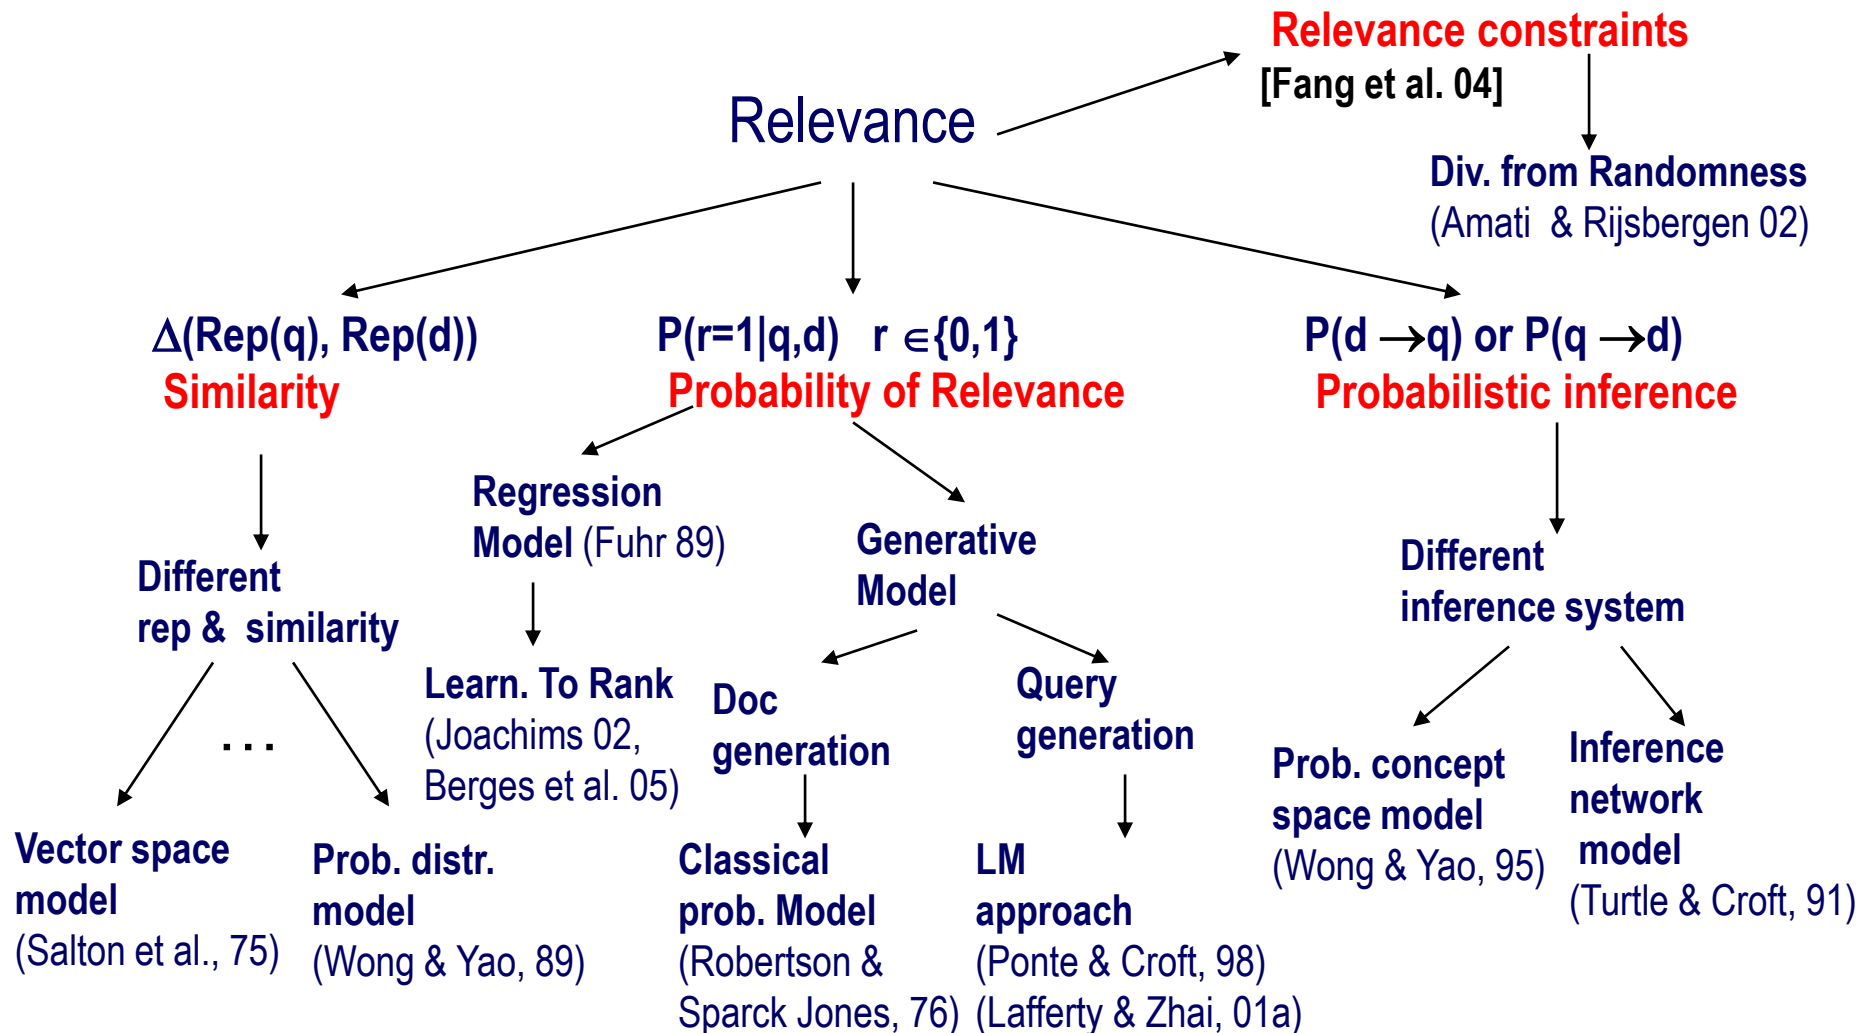
Bayes risk for choice (D, π)

Traditional Solution

- Set-based models (choose D) → Boolean model
- Ranking models (choose π)
 - Independent loss ← Pointwise
 - Relevance-based loss → { Probabilistic relevance model
Generative Relevance Theory
 - Distance-based loss → { Vector-space Model
Two-stage LM
KL-divergence model
 - Dependent loss ← Pairwise/Listwise
 - MMR loss → Subtopic retrieval model
 - MDR loss

Unsupervised!

Traditional Notion of Relevance



Broader Notion of Relevance

- Traditional view

- Content-driven

- Vector space model
 - Probability relevance model
 - Language model




Query-Document specific
Unsupervised

- Modern view

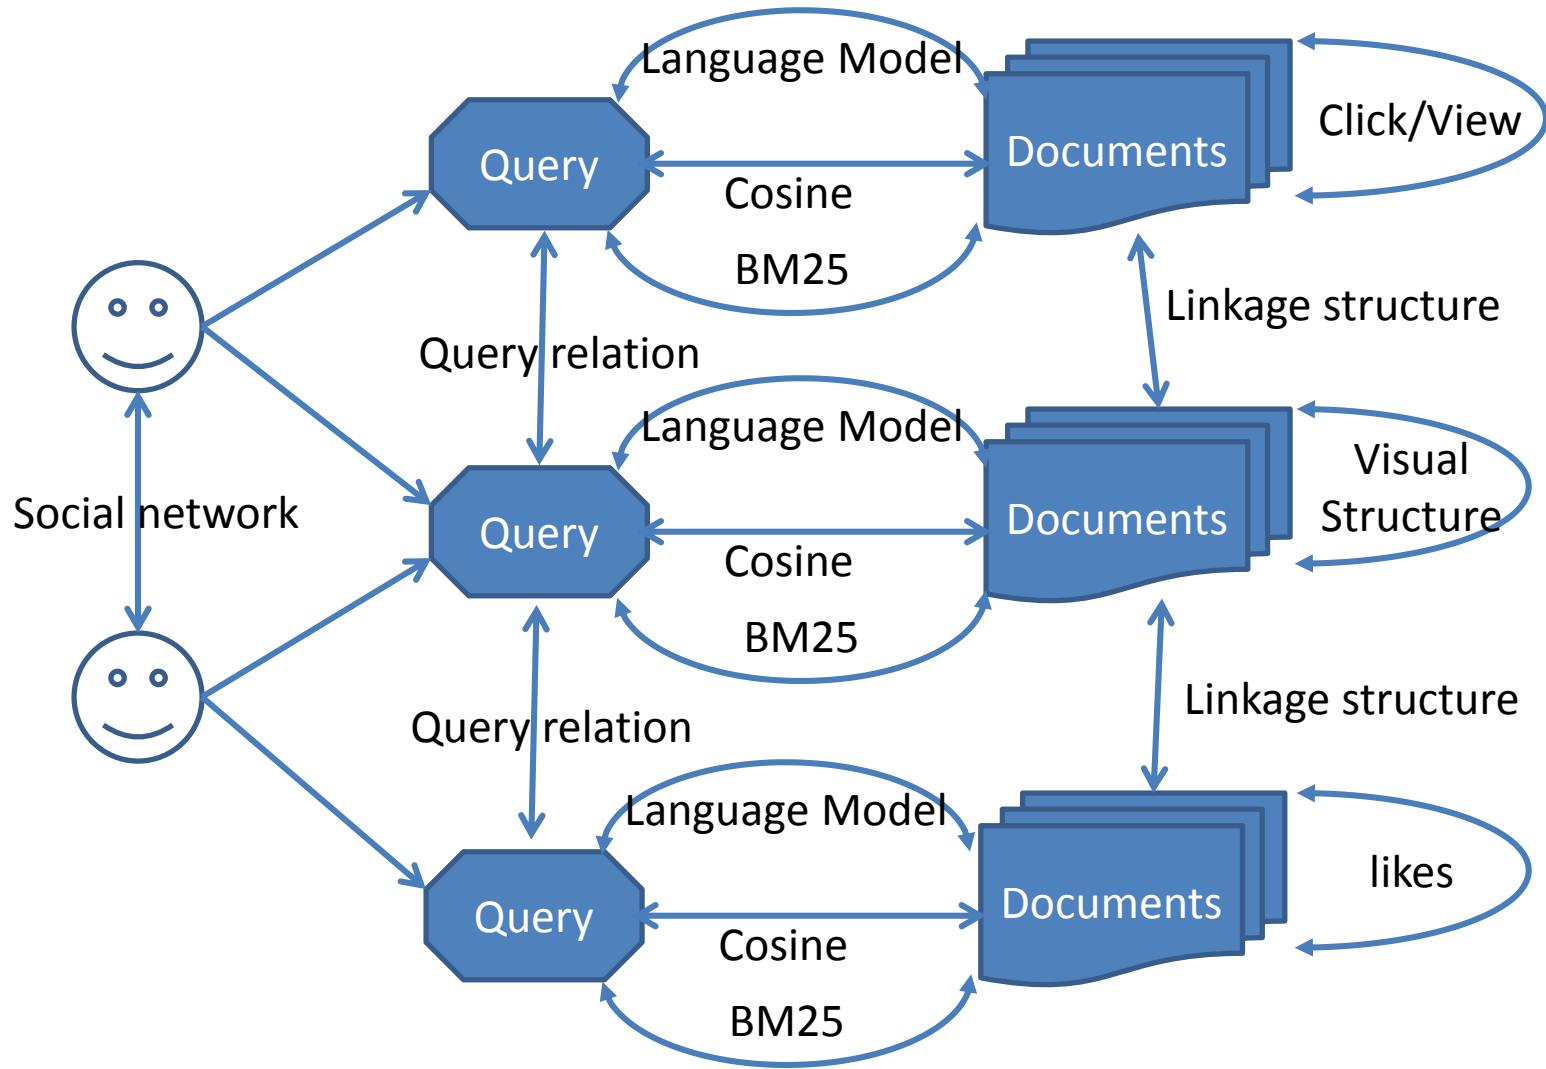
- Anything related to the quality of the document

- Clicks/views
 - Link structure
 - Visual structure
 - Social network
 -



Query, Document, Query-
Document specific
Supervised

Broader Notion of Relevance



Future

- Tight
- Fast
- Large
- Wide



Resources

- Books
 - Liu, Tie-Yan. *Learning to rank for information retrieval*. Vol. 13. Springer, 2011.
 - Li, Hang. "Learning to rank for information retrieval and natural language processing." *Synthesis Lectures on Human Language Technologies* 4.1 (2011): 1-113.
- Helpful pages
 - http://en.wikipedia.org/wiki/Learning_to_rank
- Packages
 - RankingSVM: <http://svmlight.joachims.org/>
 - RankLib: <http://people.cs.umass.edu/~vdang/ranklib.html>
- Data sets
 - LETOR <http://research.microsoft.com/en-us/um/beijing/projects/letor//>
 - Yahoo! Learning to rank challenge <http://learningtorankchallenge.yahoo.com/>

References

- Liu, Tie-Yan. "Learning to rank for information retrieval." *Foundations and Trends in Information Retrieval* 3.3 (2009): 225-331.
- Cossock, David, and Tong Zhang. "Subset ranking using regression." *Learning theory* (2006): 605-619.
- Shashua, Amnon, and Anat Levin. "Ranking with large margin principle: Two approaches." *Advances in neural information processing systems* 15 (2003): 937-944.
- Joachims, Thorsten. "Optimizing search engines using clickthrough data." *Proceedings of the eighth ACM SIGKDD*. ACM, 2002.
- Freund, Yoav, et al. "An efficient boosting algorithm for combining preferences." *The Journal of Machine Learning Research* 4 (2003): 933-969.
- Zheng, Zhaohui, et al. "A regression framework for learning ranking functions using relative relevance judgments." *Proceedings of the 30th annual international ACM SIGIR*. ACM, 2007.

References

- Joachims, Thorsten, et al. "Accurately interpreting clickthrough data as implicit feedback." Proceedings of the 28th annual international ACM SIGIR. ACM, 2005.
- Burges, C. "From ranknet to lambdarank to lambdamart: An overview." Learning 11 (2010): 23-581.
- Xu, Jun, and Hang Li. "AdaRank: a boosting algorithm for information retrieval." Proceedings of the 30th annual international ACM SIGIR. ACM, 2007.
- Yue, Yisong, et al. "A support vector method for optimizing average precision." Proceedings of the 30th annual international ACM SIGIR. ACM, 2007.
- Taylor, Michael, et al. "SoftRank: optimizing non-smooth rank metrics." Proceedings of the international conference WSDM. ACM, 2008.
- Cao, Zhe, et al. "Learning to rank: from pairwise approach to listwise approach." Proceedings of the 24th ICML. ACM, 2007.




Thank you!

- Q&A


Recap of last lecture

- Goal
 - Design the ranking module for Bing.com



Basic Search Engine Architecture

- The anatomy of a Web search engine
– Sergey Brin and Larry Page's *network*



external


puter
): 107-117.

← Your Job

Learning to Rank

- Given: (query, document) pairs represented by a set of relevance estimators, a.k.a., features

QueryID	DocID	BM25	LM	PageRank	Label
0001	0001	1.6	1.1	0.9	0
0001	0002	2.7	1.9	0.2	1

- Needed: a way of combining the estimators
 - $f(q, \{d\}_{i=1}^D) \rightarrow \text{ordered } \{d\}_{i=1}^D$
- Criterion: optimize IR metrics  **Key!**
 - P@k, MAP, NDCG, etc.

Challenge: how to optimize?

- Order is essential!
 - $f \rightarrow \mathbf{order} \rightarrow \text{metric}$
- Evaluation metrics are not continuous and not differentiable



Approximating the Objects!

- Pointwise
 - Fit the relevance labels individually
- Pairwise
 - Fit the relative orders
- Listwise
 - Fit the whole order



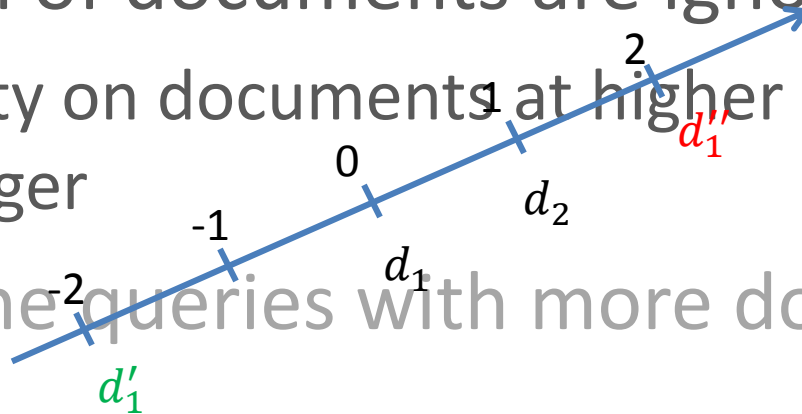
Pointwise Learning to Rank

- Ideally perfect relevance prediction leads to perfect ranking
 - $f \rightarrow \mathbf{score} \rightarrow \text{order} \rightarrow \text{metric}$
- Reduce ranking problem to
 - Regression
 - Classification




Deficiency

- Cannot directly optimize IR metrics
 - $(0 \rightarrow 1, 2 \rightarrow 0)$ worse than $(0 \rightarrow -2, 2 \rightarrow 4)$
- Position of documents are ignored
 - Penalty on documents at higher positions should be larger
- Favor the queries with more documents



Pairwise Learning to Rank

- Ideally perfect partial order leads to perfect ranking
 - $f \rightarrow$ **partial order** \rightarrow order \rightarrow metric
- Ordinal regression
 - $O(f(Q, D), Y) = \sum_{i \neq j} \delta(y_i > y_j) \delta(f(q_i, d_i) < f(q_i, d_i))$
 - Relative ordering between different documents is significant



RankingSVM

Thorsten Joachims, KDD'02


- Minimizing the number of mis-ordered pairs

linear combination of features
 $y_1 > y_2, y_2 > y_3, y_1 > y_4$
 $f(q, d) = w^T X_{q,d}$

minimize: $V(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w} \cdot \vec{w} + C \sum \xi_{i,j,k}$

subject to:

QueryID	DocID	BM25	PageRank	Label
0001	0001	1.6	0.9	4
0001	0002	2.7	0.2	2
0001	0003	1.3	0.2	1
0001	0004	1.2	0.7	0



RankingSVM

Thorsten Joachims, KDD'02

- Minimizing the number of mis-ordered pairs

$$\text{minimize:} \quad V(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w} \cdot \vec{w} + C \sum \xi_{i,j,k}$$


subject to:

$$\forall (d_i, d_j) \in r_1^* : \vec{w}\Phi(q_1, d_i) \geq \vec{w}\Phi(q_1, d_j) + 1 - \xi_{i,j,1}$$

...

$$\forall (d_i, d_j) \in r_n^* : \vec{w}\Phi(q_n, d_i) \geq \vec{w}\Phi(q_n, d_j) + 1 - \xi_{i,j,n}$$

$$\forall i \forall j \forall k : \xi_{i,j,k} \geq 0$$



General Idea of Pairwise Learning to Rank

- For any pair of $y_i > y_j$

